

Machine Learning Homework Sheet 1

k -Nearest Neighbors and Decision Trees

1 Dataset

You are free to do the following exercises completely by hand or use a computer to help speed things up (e.g., using Python). You should, however, show the basic steps of your work and implement your own “helpers” instead of blindly using existing implementations. Using a machine learning toolbox and copying the result will not help you understand the material.

The table below gives you a feature matrix \mathbf{X} together with the targets \mathbf{y} , where each row i represents one sample. This data is also available on Piazza as `01_homework_dataset.csv`. The column with names for each datapoint i can help you reference specific data points.

i	$x_{i,1}$	$x_{i,2}$	$x_{i,3}$	y_i
A	5.5	0.5	4.5	2
B	7.4	1.1	3.6	0
C	5.9	0.2	3.4	2
D	9.9	0.1	0.8	0
E	6.9	-0.1	0.6	2
F	6.8	-0.3	5.1	2
G	4.1	0.3	5.1	1
H	1.3	-0.2	1.8	1
I	4.5	0.4	2.0	0
J	0.5	0.0	2.3	1
K	5.9	-0.1	4.4	0
L	9.3	-0.2	3.2	0
M	1.0	0.1	2.8	1
N	0.4	0.1	4.3	1
O	2.7	-0.5	4.2	1

2 Decision Trees

Problem 1: Build a decision tree T for your data $\{\mathbf{X}, \mathbf{y}\}$. Consider all possible splits for all features and use the Gini index to build your tree. Build the tree only to a depth of two! Provide at least the value of the final Gini index at each node and the distribution of classes at each leaf.

Problem 2: Use the final tree T from the previous problem to classify the vectors $\mathbf{x}_a = (4.1, -0.1, 2.2)^T$ and $\mathbf{x}_b = (6.1, 0.4, 1.3)^T$. Provide both your classification \hat{y}_a and \hat{y}_b and their respective probabilities $p(c = \hat{y}_a \mid \mathbf{x}_a, T)$ and $p(c = \hat{y}_b \mid \mathbf{x}_b, T)$

Upload a single PDF file with your solution to Moodle by 29.10.2018, 9:59am CET. We recommend to typeset your solution (using L^AT_EX or Word), but handwritten solutions are also accepted.

If your handwritten solution is illegible, it won't be graded and you waive your right to dispute that.

3 k -Nearest Neighbors

Problem 3: Load the notebook `01_homework_kNN.ipynb` from piazza. Fill in the missing code and run the notebook. Convert the evaluated notebook to PDF and add it to the printout of your homework.

Note: We suggest that you use Anaconda for installing Python and Jupyter, as well as for managing packages. We recommend that you use Python 3.

For more information on Jupyter notebooks, consult the Jupyter documentation. Instructions for converting the Jupyter notebooks to PDF are provided within the notebook.

Problem 4: Classify the two vectors \mathbf{x}_a and \mathbf{x}_b given in Problem 2 with the k -nearest neighbors algorithm. Use $k = 3$ and Euclidean distance.

Problem 5: Now, consider y_i to be real-valued targets rather than classes. Perform 3-NN regression to label the vectors from Problem 2.

Problem 6: Look at the data. Which problem do you see w.r.t. building a Euclidean distance-based k -NN model on \mathbf{X} ? How can you compensate for this problem? Does this problem also arise when training a decision tree?